

Raaghav_94_NLP2_Sentiment_Analysis

January 26, 2024

1 Sentiment Analysis

1.1 1. Explain the pipeline for developing sentiment analysis task.

Sentiment Analysis is the process of determining whether a piece of text is positive, negative or neutral. The pipeline for this task is as follows:

- **Loading the Data**
 - Selecting and Downloading the appropriate dataset.
- **Data Cleaning and Preprocessing**
 - Retain only the text and their sentiments.
 - Remove Stopwords and Punctuations.
 - Lowercase the text.
 - Tokenize the dataset.
 - Lemmatize the words.
- **Feature Representation**
 - Vectorize the text using BoW, TF-IDF or word embeddings like Word2Vec.
- **Model Building**
 - Classifier such as Logistic Regression, SVM, Decision Tree, Naive Bayes etc.
- **Model Evaluation**
 - A test set is preprocessed with the same techniques and input to the model to classify and evaluate the performance.

1.2 2. Perform cleaning and preprocessing of text.

```
[ ]: import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from gensim.parsing.preprocessing import remove_stopwords
from gensim.utils import simple_preprocess
from nltk.stem.wordnet import WordNetLemmatizer
```

1.2.1 Load Dataset

```
[ ]: data = pd.read_csv("archive/train.csv", encoding='unicode_escape')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27481 entries, 0 to 27480
```

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	textID	27481 non-null	object
1	text	27480 non-null	object
2	selected_text	27480 non-null	object
3	sentiment	27481 non-null	object
4	Time of Tweet	27481 non-null	object
5	Age of User	27481 non-null	object
6	Country	27481 non-null	object
7	Population -2020	27481 non-null	int64
8	Land Area (Km ²)	27481 non-null	float64
9	Density (P/Km ²)	27481 non-null	int64

dtypes: float64(1), int64(2), object(7)

memory usage: 2.1+ MB

```
[ ]: data = data[['selected_text', 'sentiment']]
data.columns = ['sentences', 'sentiment']
data.dropna(inplace=True)
data.head(3)
```

```
[ ]:
      sentences sentiment
0  I`d have responded, if I were going    neutral
1                      Sooo SAD    negative
2                bullying me    negative
```

```
[ ]: le = LabelEncoder()
data['sentiment'] = le.fit_transform(data['sentiment'])
data.head(3)
```

```
[ ]:
      sentences  sentiment
0  I`d have responded, if I were going      1
1                      Sooo SAD      0
2                bullying me      0
```

1.3 Preprocessing the Data

```
[ ]: lemma = WordNetLemmatizer()

def preprocess(text):
    text = simple_preprocess(remove_stopwords(text), deacc=True)
    return [lemma.lemmatize(str(word)) for word in text]
```

```
[ ]: data['sentences'] = data['sentences'].apply(preprocess)
sentences = data["sentences"].values.tolist()
```

```
[ ]: sentences_list = [" ".join(i) for i in sentences]
```

1.4 3.Generate representations using

```
[ ]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
     from gensim.models.word2vec import Word2Vec
```

1.4.1 Bag of Words

```
[ ]: cv = CountVectorizer()
     count_matrix = cv.fit_transform(sentences_list)
     count_array = count_matrix.toarray()
```

1.4.2 TF-IDF

```
[ ]: tfidf = TfidfVectorizer()
     tfidf_matrix = tfidf.fit_transform(sentences_list)
     tfidf_array = tfidf_matrix.toarray()
```

1.4.3 Word2Vec - Continuous Bag of Words

```
[ ]: cbow = Word2Vec(sentences, vector_size=100, window=5, min_count=2, sg=0)
```

```
[ ]: vocab = cbow.wv.index_to_key

def get_mean_vector(model, sentence):
    words = [word for word in sentence if word in vocab]
    if len(words) >= 1:
        return np.mean(model.wv[words], axis=0)
    return np.zeros((100,))
```

```
[ ]: cbow_array = []
     for sentence in sentences:
         mean_vec = get_mean_vector(cbow, sentence)
         cbow_array.append(mean_vec)

     cbow_array = np.array(cbow_array)
```

1.4.4 Word2Vec - Skip-gram

```
[ ]: sg = Word2Vec(sentences, vector_size=100, window=5, min_count=2, sg=1)
```

```
[ ]: vocab = sg.wv.index_to_key

def get_mean_vector(model, sentence):
    words = [word for word in sentence if word in vocab]
    if len(words) >= 1:
        return np.mean(model.wv[words], axis=0)
    return np.zeros((100,))
```

```
[ ]: sg_array = []
    for sentence in sentences:
        sg_array.append(get_mean_vector(sg, sentence))

sg_array = np.array(sg_array)
```

1.5 4. Classify the data using appropriate machine learning techniques to generate labels.

```
[ ]: from sklearn.naive_bayes import MultinomialNB
    from sklearn.tree import DecisionTreeClassifier
    from sklearn.metrics import classification_report

y_train = data['sentiment'].values
```

```
[ ]: bow_model = MultinomialNB()
    bow_model.fit(count_array, y_train)
    bow_model.score(count_array, y_train)
```

```
[ ]: 0.8568413391557497
```

```
[ ]: tfidf_model = MultinomialNB()
    tfidf_model.fit(tfidf_array, y_train)
    tfidf_model.score(tfidf_array, y_train)
```

```
[ ]: 0.8592430858806405
```

```
[ ]: cbow_model = DecisionTreeClassifier()
    cbow_model.fit(cbow_array, y_train)
    cbow_model.score(cbow_array, y_train)
```

```
[ ]: 0.9702328966521107
```

```
[ ]: sg_model = DecisionTreeClassifier()
    sg_model.fit(sg_array, y_train)
    sg_model.score(sg_array, y_train)
```

```
[ ]: 0.9702328966521107
```

1.6 5. Analyze the labels and explain the impact of embedding techniques in misclassification.

```
[ ]: x_test = ["What is not to like about this product.",
    "Not bad.",
    "Not an issue.",
    "Not buggy.",
    "Not happy.",
    "Not user-friendly.",
```

```

"Not good.",
"Is it any good?",
"I do not dislike horror movies.",
"Disliking horror movies is not uncommon.",
"Sometimes I really hate the show.",
"I love having to wait two months for the next series to come out!",
"The final episode was surprising with a terrible twist at the end.",
"The film was easy to watch but I would not recommend it to my friends.",
"I LOL'd at the end of the cake scene."]

y_test = [2, 1, 1, 1, 0, 0, 0, 1, 2, 0, 0, 2, 0, 0, 1]
sentiment = {0:"Negative", 1:"Neutral", 2:"Positive"}

```

```

[ ]: bow_test = [' '.join(preprocess(sentence)) for sentence in x_test]
bow_test = cv.transform(bow_test).toarray()

tfidf_test = [' '.join(preprocess(sentence)) for sentence in x_test]
tfidf_test = tfidf.transform(tfidf_test).toarray()

cbow_test_embeds = [(preprocess(sentence)) for sentence in x_test]
cbow_test = []
for sentence in cbow_test_embeds:
    cbow_test.append(get_mean_vector(cbow, sentence))
cbow_test = np.array(cbow_test)

sg_test_embeds = [(preprocess(sentence)) for sentence in x_test]
sg_test = []
for sentence in sg_test_embeds:
    sg_test.append(get_mean_vector(sg, sentence))
sg_test = np.array(sg_test)

```

```

[ ]: y_bow = bow_model.predict(bow_test)
print([sentiment[i] for i in y_bow])
print(classification_report(y_test, y_bow))

```

```

['Neutral', 'Negative', 'Negative', 'Neutral', 'Positive', 'Negative',
'Negative', 'Positive', 'Neutral', 'Negative', 'Negative', 'Neutral', 'Neutral',
'Neutral', 'Neutral']

```

	precision	recall	f1-score	support
0	0.67	0.57	0.62	7
1	0.29	0.40	0.33	5
2	0.00	0.00	0.00	3
accuracy			0.40	15
macro avg	0.32	0.32	0.32	15
weighted avg	0.41	0.40	0.40	15

```
[ ]: y_tfidf = tfidf_model.predict(tfidf_test)
print([sentiment[i] for i in y_tfidf])
print(classification_report(y_test, y_tfidf))
```

```
['Neutral', 'Negative', 'Negative', 'Neutral', 'Positive', 'Neutral',
'Negative', 'Neutral', 'Neutral', 'Negative', 'Negative', 'Neutral', 'Neutral',
'Neutral', 'Neutral']
```

	precision	recall	f1-score	support
0	0.60	0.43	0.50	7
1	0.33	0.60	0.43	5
2	0.00	0.00	0.00	3
accuracy			0.40	15
macro avg	0.31	0.34	0.31	15
weighted avg	0.39	0.40	0.38	15

```
[ ]: y_cbow = cbow_model.predict(cbow_test)
print([sentiment[i] for i in y_cbow])
print(classification_report(y_test, y_cbow))
```

```
['Neutral', 'Negative', 'Positive', 'Negative', 'Negative', 'Negative',
'Negative', 'Negative', 'Positive', 'Negative', 'Neutral', 'Neutral', 'Neutral',
'Neutral', 'Negative']
```

	precision	recall	f1-score	support
0	0.50	0.57	0.53	7
1	0.00	0.00	0.00	5
2	0.50	0.33	0.40	3
accuracy			0.33	15
macro avg	0.33	0.30	0.31	15
weighted avg	0.33	0.33	0.33	15

```
[ ]: y_sg = sg_model.predict(sg_test)
print([sentiment[i] for i in y_sg])
print(classification_report(y_test, y_sg))
```

```
['Negative', 'Negative', 'Negative', 'Negative', 'Negative', 'Negative',
'Negative', 'Positive', 'Neutral', 'Neutral', 'Negative', 'Neutral', 'Positive',
'Neutral', 'Neutral']
```

	precision	recall	f1-score	support
0	0.50	0.57	0.53	7
1	0.20	0.20	0.20	5
2	0.00	0.00	0.00	3

accuracy			0.33	15
macro avg	0.23	0.26	0.24	15
weighted avg	0.30	0.33	0.32	15

1.7 6. Discuss the limitations of each embedding technique and explain the techniques that rectify it.

- BOW
 - Sparse Representation.
 - Word Order is not considered.
 - Does not capture semantic meaning of the text.
 - Computationally Intensive.
- TF-IDF
 - Sparse Representation.
 - Does not capture semantic meaning of the text.
 - Computationally Intensive.
- Word2Vec
 - Semantic and Dense Representation.
 - CBow: Faster and Better Representation for Frequent Words.
 - SkipGram: Works well with small amount of data and Represent Rare words well.

All these still face Out-Of-Vocabulary (OOV) problem that can be resolved by using FastText which uses N-grams of the words.