

Raaghav_94_NLP4_Hate

February 14, 2024

1 Hate Speech Identification

1.1 Pipeline

- Loading the input data.
- Cleaning and Preprocessing the data.
 - Lowercasing the text data.
 - Removing punctuations and stopwords.
 - Tokenizing the sentences.
- Generate representations for the features.
 - Word2Vec, FastText
 - * Get the mean vectors for the sentences.
 - CNN, RNN
- Feeding the representation into a classifier model.
- Evaluating the model using test data with accuracy as the metric.

1.2 Import Libraries

```
[ ]: import numpy as np
import pandas as pd
import regex as re
from gensim.utils import simple_preprocess
from gensim.parsing.preprocessing import remove_stopwords
from nltk.stem.wordnet import WordNetLemmatizer
```

1.3 Load Data

```
[ ]: tweet = pd.read_csv("labeled_data.csv")
tweet = tweet[['class', 'tweet']]
tweet.head()
```

```
[ ]:      class      tweet
0         2  !!! RT @mayasolovely: As a woman you shouldn't...
1         1  !!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2         1  !!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3         1  !!!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4         1  !!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...
```

```
[ ]: X = tweet['tweet']
      y = tweet['class']

      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
      ↪random_state=0, stratify=y)
```

1.4 Preprocess Data

```
[ ]: import nltk
      nltk.download('wordnet')
```

[nltk_data] Downloading package wordnet to /root/nltk_data...

[nltk_data] Package wordnet is already up-to-date!

```
[ ]: True
```

```
[ ]: lemma = WordNetLemmatizer()

      def preprocess(text):
          text = re.sub('@[\w]+', '', text)
          text = simple_preprocess(remove_stopwords(text))
          return [lemma.lemmatize(str(word)) for word in text if word != 'rt']

      train_sentences = [preprocess(text) for text in X_train]
      test_sentences = [preprocess(text) for text in X_test]
```

1.5 Word Embeddings

```
[ ]: from gensim.models.word2vec import Word2Vec
      from gensim.models import FastText
```

1.5.1 Word2Vec

```
[ ]: cbow_train = Word2Vec(train_sentences, vector_size=100, window=5, min_count=2,
      ↪sg=0)
      cbow_test = Word2Vec(test_sentences, vector_size=100, window=5, min_count=2,
      ↪sg=0)
```

```
[ ]: train_vocab = cbow_train.wv.index_to_key
      test_vocab = cbow_test.wv.index_to_key

      def get_mean_vector(model, sentence, vocab):
          words = [word for word in sentence if word in vocab]
          if len(words) >= 1:
              return np.mean(model.wv[words], axis=0)
          return np.zeros((100,))
```

```
[ ]: cbow_array_train = []
      cbow_array_test = []

      for sentence in train_sentences:
          mean_vec = get_mean_vector(cbow_train, sentence, train_vocab)
          cbow_array_train.append(mean_vec)

      for sentence in test_sentences:
          mean_vec = get_mean_vector(cbow_test, sentence, test_vocab)
          cbow_array_test.append(mean_vec)

      cbow_array_train = np.array(cbow_array_train)
      cbow_array_test = np.array(cbow_array_test)
```

1.5.2 FastText

```
[ ]: fasttext_train = FastText(train_sentences, sg=1, workers=4, vector_size=100,
    ↪min_count=2,window=5)
      fasttext_test = FastText(test_sentences, sg=1, workers=4, vector_size=100,
    ↪min_count=2,window=5)
```

```
[ ]: fasttext_array_train = []
      fasttext_array_test = []

      for sentence in train_sentences:
          mean_vec = get_mean_vector(fasttext_train, sentence, train_vocab)
          fasttext_array_train.append(mean_vec)

      for sentence in test_sentences:
          mean_vec = get_mean_vector(fasttext_test, sentence, test_vocab)
          fasttext_array_test.append(mean_vec)

      fasttext_array_train = np.array(fasttext_array_train)
      fasttext_array_test = np.array(fasttext_array_test)
```

1.6 Model Building

```
[ ]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score
```

```
[ ]: cbow_tree = SVC()
      cbow_tree.fit(cbow_array_train, y_train)
```

```
[ ]: SVC()
```

```
[ ]: cbow_pred = cbow_tree.predict(cbow_array_test)
print(f"Accuracy for Word2Vec: {accuracy_score(y_test, cbow_pred):.4f}")
```

Accuracy for Word2Vec: 0.7660

```
[ ]: fasttext_tree = SVC()
fasttext_tree.fit(fasttext_array_train, y_train)
```

```
[ ]: SVC()
```

```
[ ]: fasttext_pred = fasttext_tree.predict(fasttext_array_test)
print(f"Accuracy for FastText: {accuracy_score(y_test, fasttext_pred):.4f}")
```

Accuracy for FastText: 0.7751

1.7 CNN

```
[ ]: from keras.models import Sequential
from keras.layers import Dense, Flatten, Conv1D, MaxPooling1D, Embedding,
↳ Bidirectional, LSTM
from keras.preprocessing import sequence
from keras.preprocessing.text import Tokenizer
```

```
[ ]: top_words = 10000

tokenizer = Tokenizer()
train_sequences = tokenizer.texts_to_sequences(train_sentences)
train_pad_sequences = sequence.pad_sequences(train_sequences, maxlen=100)

test_sequences = tokenizer.texts_to_sequences(test_sentences)
test_pad_sequences = sequence.pad_sequences(test_sequences, maxlen=100)
```

```
[ ]: model = Sequential([
    Embedding(top_words, 32, input_length=100),
    Conv1D(32, 3, padding='same', activation='relu'),
    MaxPooling1D(),
    Flatten(),
    Dense(250, activation='relu'),
    Dense(1, activation='softmax')
])

model.compile(loss='binary_crossentropy', optimizer='adam',
↳ metrics=['accuracy'])
model.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		

embedding_8 (Embedding)	(None, 100, 32)	320000
conv1d_4 (Conv1D)	(None, 100, 32)	3104
max_pooling1d_4 (MaxPoolin g1D)	(None, 50, 32)	0
flatten_4 (Flatten)	(None, 1600)	0
dense_12 (Dense)	(None, 250)	400250
dense_13 (Dense)	(None, 1)	251

```
=====
Total params: 723605 (2.76 MB)
Trainable params: 723605 (2.76 MB)
Non-trainable params: 0 (0.00 Byte)
-----
```

```
[ ]: model.fit(train_pad_sequences, y_train, validation_data=(test_pad_sequences,
↪y_test), epochs=2, batch_size=128, verbose=2)
```

Epoch 1/2

155/155 - 7s - loss: -1.0494e+03 - accuracy: 0.7743 - val_loss: -5.8186e+03 -
val_accuracy: 0.7743 - 7s/epoch - 42ms/step

Epoch 2/2

155/155 - 6s - loss: -4.9316e+04 - accuracy: 0.7743 - val_loss: -1.3855e+05 -
val_accuracy: 0.7743 - 6s/epoch - 37ms/step

```
[ ]: <keras.src.callbacks.History at 0x7b6358a3f910>
```

```
[ ]: scores = model.evaluate(test_pad_sequences, y_test)
print(f"Accuracy for CNN: {(scores[1]*100):.4f}")
```

155/155 [=====] - 1s 7ms/step - loss: -138546.8906 -
accuracy: 0.7743

Accuracy for CNN: 77.4259

1.8 RNN

```
[ ]: rnn = Sequential([
    Embedding(top_words, 64, input_length=100),
    Bidirectional(LSTM(64)),
    Dense(1, activation='softmax')
])
rnn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
[ ]: history = rnn.fit(train_pad_sequences,
    ↪y_train, validation_data=(test_pad_sequences, y_test), epochs=2,
    ↪batch_size=128)
```

Epoch 1/2

155/155 [=====] - 60s 345ms/step - loss: -2.1117 -
accuracy: 0.7743 - val_loss: -3.7022 - val_accuracy: 0.7743

Epoch 2/2

155/155 [=====] - 53s 343ms/step - loss: -4.8768 -
accuracy: 0.7743 - val_loss: -6.0184 - val_accuracy: 0.7743

```
[ ]: scores = rnn.evaluate(test_pad_sequences, y_test)
    print(f"Accuracy for RNN: {(scores[1]*100):.4f}")
```

155/155 [=====] - 7s 47ms/step - loss: -6.0184 -
accuracy: 0.7743

Accuracy for RNN: 77.4259

1.9 Result

```
[ ]: from sklearn.metrics import classification_report
    print(classification_report(y_test, cbow_pred, zero_division=True))
```

	precision	recall	f1-score	support
0	1.00	0.00	0.00	286
1	0.84	0.88	0.86	3838
2	0.43	0.49	0.46	833
accuracy			0.77	4957
macro avg	0.76	0.46	0.44	4957
weighted avg	0.78	0.77	0.75	4957

- Predicting the ‘neutral’ class is very difficult for the models.
- It classifies all the tweets only as either positive or negative.